

Thanks so much for applying to join us as a Senior AI Builder, we're excited to learn more about you!

As the next step, we'd love for you to complete a short qualification project. This role is hands-on from day one, building production AI agents on the Anthropic stack for real clients, so the best way for us to see how you think and ship is to actually see you build something small.

Deadline: Please send us your submission within 3 days of receiving this technical test.

What we'd like you to build

A working AI agent that answers questions over a public documentation site of your choice (pick something you find interesting, a product's docs, an open-source project, a company knowledge base, whatever). The agent should:

1. Ingest the site's content into a vector store (we strongly prefer Supabase + pgvector since that's our production stack, but if you can't ship it in time, use what you know and note the tradeoff).
2. Answer questions over that content with proper citations back to the source.
3. Be exposed via a small FastAPI endpoint, so we can see your Python.
4. Include some kind of evaluation, even just 10 to 15 question/answer pairs with a simple script that measures retrieval or answer quality (precision-at-k, LLM-as-judge, RAGAS, your call). We care a lot about this. Anyone can wire up RAG, we want to see how you think about whether it actually works.

Then pick ONE of these to go deeper on (whichever you think will impress us most):

Option A - Managed Agent Q&A

A Claude Managed Agent that handles the Q&A flow as a multi-turn conversation, with proper session management and at least one custom tool definition (e.g. a search tool, a citation formatter, a source ranker). Focus on the agent architecture, not the ingestion.

Option B - AI Copilot Integration

Wrap the Q&A flow so a non-technical user can trigger it from their daily environment. This could be a Claude Cowork Skill or Plugin, a Cursor extension, a GitHub Copilot plugin, or any similar AI copilot integration you work with regularly. (Note: if you don't have Claude Cowork access, this is the easiest path to demonstrate the same concept.)

Option C – MCP Server

An MCP server that exposes the knowledge base as a tool any Claude client can connect to, with proper tool definitions, clear input/output schema, and a short demo of it working in a Claude client.

Pick the one you'd actually enjoy building. We're more interested in seeing one thing done well than three things done shallow.

Scope and time

Aim for 4 to 6 hours. We're not looking for a polished product, we're looking for signal on how you build, what tradeoffs you make, and how you handle the Anthropic-specific tooling. If something would take longer to do properly, write a short note on what you'd do differently with more time.

What to send us

- A GitHub repo public link with your code, a clear README, and any setup instructions.
- A short Loom (5 to 10 minutes) where you introduce yourself, walk us through what you built, show it working, and explain the key decisions you made. The Loom matters: this role involves talking with clients, so we want to see how you communicate technical work.
- A few sentences on what you'd build next if this were day one of the job.

Let us know if you have any questions, looking forward to seeing what you create!